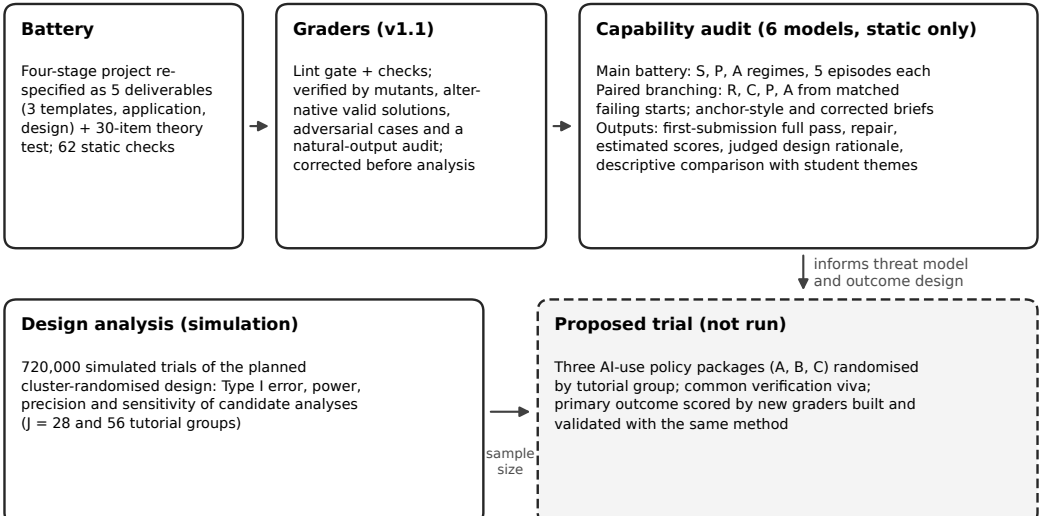
Battery
Four-stage project re-specified as 5 deliverables (3 templates, application, design) + 30-item theory test; 62 static checks
Graders (v1.1)
Lint gate + checks; verified by mutants, alternative valid solutions, adversarial cases and a natural-output audit; corrected before analysis
Capability audit (6 models, static only)
Main battery: S, P, A regimes, 5 episodes each
Paired branching: R, C, P, A from matched failing starts; anchor-style and corrected briefs
Outputs: first-submission full pass, repair, estimated scores, judged design rationale, descriptive comparison with student themes
informs threat model and outcome design
Design analysis (simulation)
720,000 simulated trials of the planned cluster-randomised design: Type I error, power, precision and sensitivity of candidate analyses (J = 28 and 56 tutorial groups)
sample size
Proposed trial (not run)
Three AI-use policy packages (A, B, C) randomised by tutorial group; common verification viva; primary outcome scored by new graders built and validated with the same method